

# **OpenCL Imaging on The GPU: Optical Flow**

**James Fung**  
**Developer Technology, NVIDIA**

# **Pyramidal Lucas Kanade Optical Flow in OpenCL**

Algorithm by Jean-Yves Bouguet (Intel)  
Implementation by James Fung (NVIDIA)


# Optical Flow

- **Calculate movement of selected points in pairs of images**
- **Applications:**
  - Image stabilization
  - Feature tracking
  - Video encoding
- **May be used to track**
  - a few select points: sparse optical flow
  - All image points: dense optical flow
    - Computationally intensive!





Pyramidal Lucas Kanade Optical Flow in OpenCL



# Algorithm

- For each point  $u$  on image  $I$ , find the corresponding point  $v$  on image  $J$
- Calculate spatial gradient  $G$  in vicinity of each point  $u$
- Iteratively solve for flow:

Calculate image mismatch vector  $\mathbf{b}$

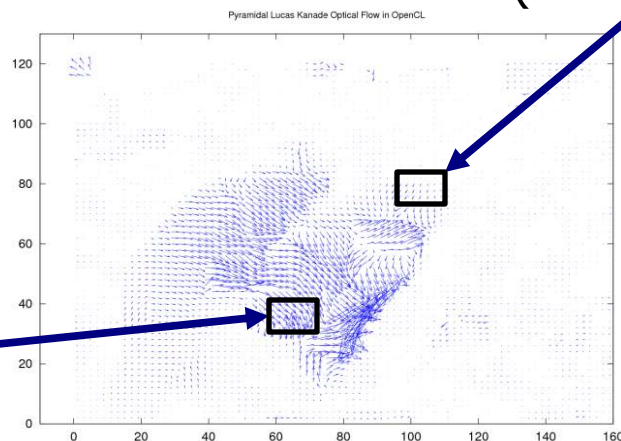
Estimate motion  $\mathbf{v} = \mathbf{G}^{-1} \mathbf{b}(x', y')$

Update position  $(x', y') = (x' + v_x, y' + v_y)$

Repeat for  $N$  iterations or until convergence

Large Motion  
(3-7 pixels)

Small motion  
( $< 3$  pixels)



# Algorithm

*Downsample*

*Scharr Edge filter*

*Solve for flow*  
 $v = G^{-1}b$

**Source image pair**


*Precompute  $G$  from  $I$*

Image I


Image J

$G_{2 \times 2}$

Flow ( $v$ )



**Flow Result**  
 $v = G^{-1}b$   
*Iterate*



*Iterate  $v = G^{-1}b$*



*Iterate  $v = G^{-1}b$*

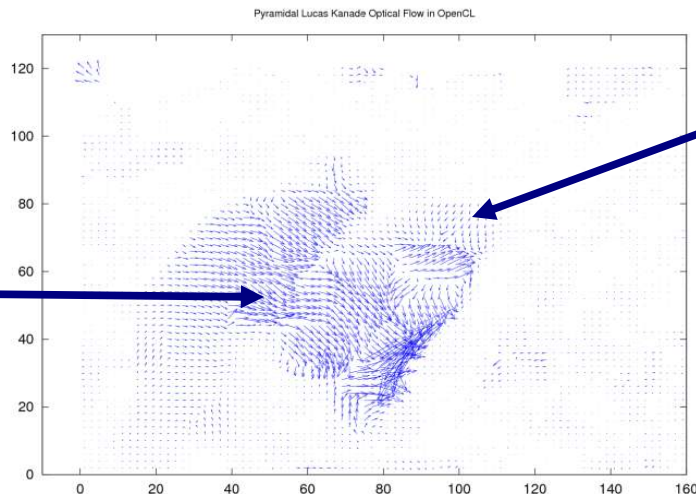


# Optical Flow in OpenCL

- **Embarrassingly parallel**
- **Computationally intensive**
- **Can take advantage of special-purpose GPU hardware:**
  - Texture cache for 2D spatial locality
  - Hardware bilinear interpolation for sub-pixel lookups
  - Local Memory
- **Can use OpenCL-OpenGL interoperability for visualization**

# Texture Cache

- **Different areas of image have different amounts of motion, but motion is spatially coherent**
  - Lookup window offset varies inside the image
  - *Texture cache* captures 2D locality



Large Motion  
(3-7 pixels)

Small motion  
( $< 3$  pixels)

# Hardware Interpolation

- Sub-pixel accuracy & sampling is crucial
- Between iterations,  $x, y$  is non-integer


$$\partial I_k(x, y) = I^L(x, y) - J^L(x + g_x^L + v_x^L, y + g_y^L + v_y^L)$$

$$b_k = \sum_{x=p_x-w_x}^{p_x+w_x} \sum_{y=p_y-w_y}^{p_y+w_y} \begin{bmatrix} \partial I_k(x, y) I_x(x, y) \\ \partial I_k(x, y) I_y(x, y) \end{bmatrix}$$

Image level  $L$   
Pixel center  $p$   
Window  $w$   
Guess  $g$  from previous level  $L+1$

- Use texture hardware linear interpolation


```
sampler_t bilinSample =  
    ... | CLK_FILTER_LINEAR;  
...  
float Jsample = read_imagef( J_float,  
    bilinSampler, Jidx+(float2)(i,j) ).x;  
...
```



# Using OpenCL Local Memory

- **Lookup locations for I, Ix and Iy are static per iteration**
- **Can be explicitly cached in Local Memory**
  - Local memory is on-chip
  - Local memory is faster to access, useful for repeated access

# Effect of moving from Texture to Local Memory Usage



Middlebury "Minicooper" data set  
Frames: 10-11 Resolution: 640 x 480 Greyscale  
Convergence: 0.0004px  
Flow compute time for 3 pyramid levels shown

# Local/Texture Buffers


# Optical Flow Demo

- **Pyramidal Lucas Kanade Optical Flow**
- **Visualization done on GPU by sharing data between OpenGL & OpenCL**



# Optical Flow Performance

## LK Pyramidal Optical Flow GTX 460 & GTX 580




Middlebury "Minicooper" data set  
Frames: 10-11 Resolution: 640x480 Greyscale  
Convergence: 0.0004px  
Pyramid 3-level flow compute time shown

# OpenCL Performance Analysis with the CUDA Visual Profiler 4.0

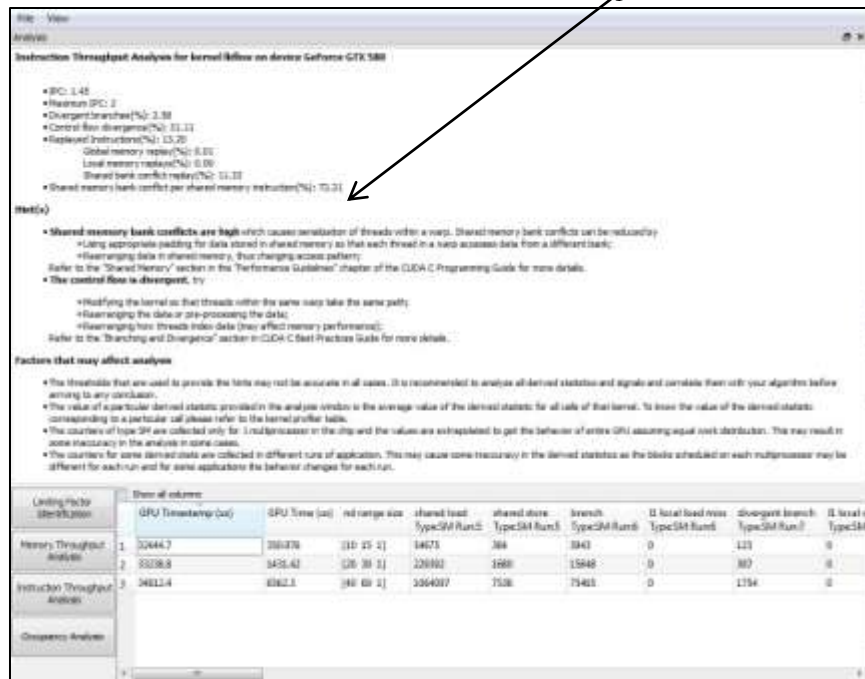
Local memory replays(%): 0.00  
Shared bank conflict replay(%): 11.33  
• Shared memory bank conflict per shared memory instruction(%): 72.21

(s)

- Shared memory bank conflicts are high which causes serialization of
  - Using appropriate padding for data stored in shared memory so that
  - Rearranging data in shared memory, thus changing access pattern



texture hit rate %
86.1799
97.5473
96.0951
82.5859
93.3682





# Thank You!

## Contact:

James Fung  
jfung@nvidia.com